## Al Movie Character Chatbot -Internship Ladder Challenge 🚀



### Goal: Progressively Build & Scale a Backend Chatbot

This internship challenge is **designed** as a ladder, where you start with a simple implementation and progressively tackle more complex, real-world challenges.

Each **milestone** gets harder, and you **earn points** for each level completed.

### Scoring System

🏅 Level	⊚ Task	Points
Level 1	Basic API Chatbot	10
Level 2	Store & Retrieve Movie Script Data	20
Level 3	Implement RAG with Vector Search	30
Level 4	Scale System to Handle High Traffic	40
Level 5	Optimise for Latency & Deploy	50
BONUS	Fine-Tune AI / Add Voice / WebSockets	20-50

🏆 Top Performers: If you score 100+ points, you are among the top candidates and may receive direct internship offers.

### 📚 Movie Script Sources

For this project, you may find useful movie scripts at the following websites:

- 1. IMSDb (Internet Movie Script Database) https://www.imsdb.com/
- 2. SimplyScripts <a href="https://www.simplyscripts.com/">https://www.simplyscripts.com/</a>
- 3. The Script Lab https://thescriptlab.com/

### Submission Guidelines

### Where to Submit:

- Candidates should submit their work via GitHub.
- Provide a **GitHub repo link** containing:
  - Source code
  - o README file
  - API Documentation

#### What to Include in the Submission:

- **README.md** explaining how to set up and run the project.
- API documentation (Postman collection, OpenAPI schema, or Markdown file).
- Screenshots / Demo Video (if applicable).

### Level 1: Basic API Chatbot (10 Points)

✓ Task: Build a simple backend that allows users to chat with a movie character.

### **Requirements:**

- Use OpenAl GPT API (or an open-source LLM like Llama2).
- Create a simple REST API:
  - POST /chat Accepts {character, user\_message} and returns a response.
- Responses should mimic the character's personality (even if manually defined).
- **© Goal:** Show basic API integration & response handling.
- **Expected Time:** ~2-3 hours

# ✓ Level 2: Store & Retrieve Movie Script Data (20 Points)

✓ Task: Instead of relying solely on the Al model, store real movie dialogues and use them for responses.

### Requirements:

- Scrape or download movie scripts from
- Store dialogues per character in a database (PostgreSQL, MongoDB, or SQLite).
- Modify POST /chat to:
  - First, search the database for an exact or closely matching line.
  - o If found, return that line.
  - If not found, generate a response using AI.
- **©** Goal: Improve realism by retrieving actual character dialogues instead of generating everything from scratch.
- **Expected Time:** ~5-7 hours

## ✓ Level 3: Implement RAG with Vector Search (30 Points)

✓ Task: Improve chat accuracy using Retrieval-Augmented Generation (RAG) with semantic search.

### Requirements:

- Convert movie scripts into vector embeddings (use OpenAl Embeddings API, FAISS, Pinecone, or ChromaDB).
- Store embeddings in a vector database.
- Modify POST /chat:
  - Retrieve the most relevant dialogue using a similarity search.
  - Pass the retrieved dialogue as context to the AI model before generating a response.
- @ Goal: Make Al more accurate by giving it real dialogues before generating responses.
- Expected Time: ~10-12 hours

## **1** Level 4: Scale System to Handle High Traffic (40 Points)

✓ Task: Optimize the backend to handle 100,000+ API requests per second.

#### **Requirements:**

- Add caching (Redis) to speed up dialogue retrieval.
- Implement rate limiting (5 requests/sec per user) to prevent abuse.
- Use async programming (Celery, Task Queues, or FastAPI's async) to improve performance.
- Run load tests (using Locust or K6) to benchmark performance.
- @ Goal: Ensure the chatbot can handle large user loads without breaking.
- Expected Time: ~12-15 hours

### Level 5: Optimize for Latency & Deploy (50 Points)

Task: Reduce response times to under 500ms and deploy the chatbot.

### Requirements:

- Implement WebSockets for real-time chat instead of REST APIs.
- Deploy the backend on AWS, DigitalOcean, or Vercel.
- Add monitoring tools (Prometheus + Grafana) to track performance.
- Store chat history per user in a database.
- @ Goal: Make the chatbot fast, scalable, and production-ready.
- **Expected Time:** ~15-20 hours

### Frequently Asked Questions:

Will keep updating FAQs as questions come.

Q: How to ask questions?

A: You can ask by commenting on this document or by asking in the internshala chat.

Q: Do we need to follow the Levels sequentially? A: No, Build in order you like.